

# [v13,02/12] perf/x86/intel/pt: Export pt\_cap\_get()

**Message ID** 1540368316-12998-3-git-send-email-luwei.kang@intel.com

**State** New

**Headers** show

**Series** Intel Processor Trace virtualization enabling

**Related** show

10654359

[diff \(/patch/10654359/raw/\)](/patch/10654359/raw/)[mbox \(/patch/10654359/mbox/\)](/patch/10654359/mbox/)[series \(/series/34463/mbox/\)](/series/34463/mbox/)

## Commit Message

Kang, Luwei (/project/kvm/list/?submitter=168537)

Oct. 24, 2018, 8:05 a.m. UTC

**From:** Chao Peng <chao.p.peng@linux.intel.com>

pt\_cap\_get() is required by the upcoming PT support in KVM guests.

Export it and move the capabilities enum to a global header.

As a global functions, "pt\_\*" is already used for ptrace and other things, so it makes sense to use "intel\_pt\_\*" as a prefix.

Acked-by: Song Liu &lt;songliubraving@fb.com&gt;

**Signed-off-by:** Chao Peng <chao.p.peng@linux.intel.com>**Signed-off-by:** Luwei Kang <luwei.kang@intel.com>

---

```

arch/x86/events/intel/pt.c      | 49 ++++++-----
arch/x86/events/intel/pt.h      | 21 -----
arch/x86/include/asm/intel_pt.h | 23 ++++++
3 files changed, 49 insertions(+), 44 deletions(-)

```

## Patch

10654359

[diff \(/patch/10654359/raw/\)](/patch/10654359/raw/)[mbox \(/patch/10654359/mbox/\)](/patch/10654359/mbox/)[series \(/series/34463/mbox/\)](/series/34463/mbox/)

```

diff --git a/arch/x86/events/intel/pt.c b/arch/x86/events/intel/pt.c
index 8d016ce..309bb1d 100644
--- a/arch/x86/events/intel/pt.c
+++ b/arch/x86/events/intel/pt.c
@@ -75,7 +75,7 @@
     PT_CAP(psb_periods,          1, CPUID_EBX, 0xffff0000),
 };

-static u32 pt_cap_get(enum pt_capabilities cap)
+u32 intel_pt_validate_hw_cap(enum pt_capabilities cap)
 {
     struct pt_cap_desc *cd = &pt_caps[cap];
     u32 c = pt_pmu.caps[cd->leaf * PT_CPUID_REGS_NUM + cd->reg];
@@ -83,6 +83,7 @@
     static u32 pt_cap_get(enum pt_capabilities cap)

     return (c & cd->mask) >> shift;
 }
+EXPORT_SYMBOL_GPL(intel_pt_validate_hw_cap);

static ssize_t pt_cap_show(struct device *cdev,
                          struct device_attribute *attr,
@@ -92,7 +93,7 @@
static ssize_t pt_cap_show(struct device *cdev,
                          container_of(attr, struct dev_ext_attribute, attr);
enum pt_capabilities cap = (long)ea->var;

-    return snprintf(buf, PAGE_SIZE, "%x\n", pt_cap_get(cap));
+    return snprintf(buf, PAGE_SIZE, "%x\n", intel_pt_validate_hw_cap(cap));
}

static struct attribute_group pt_cap_group = {
@@ -310,16 +311,16 @@
static bool pt_event_valid(struct perf_event *event)
    return false;

    if (config & RTIT_CTL_CYC_PSB) {
-        if (!pt_cap_get(PT_CAP_psb_cyc))
+        if (!intel_pt_validate_hw_cap(PT_CAP_psb_cyc))
            return false;

-        allowed = pt_cap_get(PT_CAP_psb_periods);
+        allowed = intel_pt_validate_hw_cap(PT_CAP_psb_periods);
        requested = (config & RTIT_CTL_PSB_FREQ) >>
            RTIT_CTL_PSB_FREQ_OFFSET;
        if (requested && !(allowed & BIT(requested)))
            return false;

-        allowed = pt_cap_get(PT_CAP_cycle_thresholds);
+        allowed = intel_pt_validate_hw_cap(PT_CAP_cycle_thresholds);
        requested = (config & RTIT_CTL_CYC_THRESH) >>
            RTIT_CTL_CYC_THRESH_OFFSET;
        if (requested && !(allowed & BIT(requested)))

```

```

@@ -334,10 +335,10 @@ static bool pt_event_valid(struct perf_event *event)
    * Spec says that setting mtc period bits while mtc bit in
    * CPUID is 0 will #GP, so better safe than sorry.
    */
-    if (!pt_cap_get(PT_CAP_mtc))
+    if (!intel_pt_validate_hw_cap(PT_CAP_mtc))
        return false;

-    allowed = pt_cap_get(PT_CAP_mtc_periods);
+    allowed = intel_pt_validate_hw_cap(PT_CAP_mtc_periods);
    if (!allowed)
        return false;

@@ -349,11 +350,11 @@ static bool pt_event_valid(struct perf_event *event)
}

    if (config & RTIT_CTL_PWR_EVT_EN &&
-        !pt_cap_get(PT_CAP_power_event_trace))
+        !intel_pt_validate_hw_cap(PT_CAP_power_event_trace))
        return false;

    if (config & RTIT_CTL_PTW) {
-        if (!pt_cap_get(PT_CAP_ptwrite))
+        if (!intel_pt_validate_hw_cap(PT_CAP_ptwrite))
            return false;

        /* FUPonPTW without PTW doesn't make sense */
@@ -598,7 +599,7 @@ static struct topa *topa_alloc(int cpu, gfp_t gfp)
    * In case of single-entry ToPA, always put the self-referencing END
    * link as the 2nd entry in the table
    */
-    if (!pt_cap_get(PT_CAP_topa_multiple_entries)) {
+    if (!intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries)) {
        TOPA_ENTRY(topa, 1)->base = topa->phys >> TOPA_SHIFT;
        TOPA_ENTRY(topa, 1)->end = 1;
    }

@@ -638,7 +639,7 @@ static void topa_insert_table(struct pt_buffer *buf, struct topa *topa)
    topa->offset = last->offset + last->size;
    buf->last = topa;

-    if (!pt_cap_get(PT_CAP_topa_multiple_entries))
+    if (!intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries))
        return;

    BUG_ON(last->last != TENTS_PER_PAGE - 1);
@@ -654,7 +655,7 @@ static void topa_insert_table(struct pt_buffer *buf, struct topa *topa)
static bool topa_table_full(struct topa *topa)
{
    /* single-entry ToPA is a special case */
-    if (!pt_cap_get(PT_CAP_topa_multiple_entries))
+    if (!intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries))

```

```

        return !!topa->last;

        return topa->last == TENTS_PER_PAGE - 1;
@@ -690,7 +691,8 @@ static int topa_insert_pages(struct pt_buffer *buf, gfp_t gfp)

        TOPA_ENTRY(topa, -1)->base = page_to_phys(p) >> TOPA_SHIFT;
        TOPA_ENTRY(topa, -1)->size = order;
-       if (!buf->snapshot && !pt_cap_get(PT_CAP_topa_multiple_entries)) {
+       if (!buf->snapshot &&
+       !intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries)) {
                TOPA_ENTRY(topa, -1)->intr = 1;
                TOPA_ENTRY(topa, -1)->stop = 1;
        }
@@ -725,7 +727,7 @@ static void pt_topa_dump(struct pt_buffer *buf)
                topa->table[i].intr ? 'I' : ' ',
                topa->table[i].stop ? 'S' : ' ',
                *(u64 *)&topa->table[i]);
-       if ((pt_cap_get(PT_CAP_topa_multiple_entries) &&
+       if ((intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries) &&
                topa->table[i].stop) ||
                topa->table[i].end)
                break;
@@ -828,7 +830,7 @@ static void pt_handle_status(struct pt *pt)
        * means we are already losing data; need to let the decoder
        * know.
        */
-       if (!pt_cap_get(PT_CAP_topa_multiple_entries) ||
+       if (!intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries) ||
                buf->output_off == sizes(TOPA_ENTRY(buf->cur, buf->cur_idx)->size)) {
                perf_aux_output_flag(&pt->handle,
                                PERF_AUX_FLAG_TRUNCATED);
@@ -840,7 +842,8 @@ static void pt_handle_status(struct pt *pt)
        * Also on single-entry ToPA implementations, interrupt will come
        * before the output reaches its output region's boundary.
        */
-       if (!pt_cap_get(PT_CAP_topa_multiple_entries) && !buf->snapshot &&
+       if (!intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries) &&
+       !buf->snapshot &&
                pt_buffer_region_size(buf) - buf->output_off <= TOPA_PMI_MARGIN) {
                void *head = pt_buffer_region(buf);

@@ -931,7 +934,7 @@ static int pt_buffer_reset_markers(struct pt_buffer *buf,

        /* single entry ToPA is handled by marking all regions STOP=1 INT=1 */
-       if (!pt_cap_get(PT_CAP_topa_multiple_entries))
+       if (!intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries))
                return 0;

        /* clear STOP and INT from current entry */
@@ -1082,7 +1085,7 @@ static int pt_buffer_init_topa(struct pt_buffer *buf, unsigned long nr_pages,

```

```

    pt_buffer_setup_topa_index(buf);

    /* link last table to the first one, unless we're double buffering */
-   if (pt_cap_get(PT_CAP_topa_multiple_entries)) {
+   if (intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries)) {
        TOPA_ENTRY(buf->last, -1)->base = buf->first->phys >> TOPA_SHIFT;
        TOPA_ENTRY(buf->last, -1)->end = 1;
    }
@@ -1153,7 +1156,7 @@ static int pt_addr_filters_init(struct perf_event *event)
    struct pt_filters *filters;
    int node = event->cpu == -1 ? -1 : cpu_to_node(event->cpu);

-   if (!pt_cap_get(PT_CAP_num_address_ranges))
+   if (!intel_pt_validate_hw_cap(PT_CAP_num_address_ranges))
        return 0;

    filters = kzalloc_node(sizeof(struct pt_filters), GFP_KERNEL, node);
@@ -1202,7 +1205,7 @@ static int pt_event_addr_filters_validate(struct list_head *filters)
    return -EINVAL;
}

-   if (++range > pt_cap_get(PT_CAP_num_address_ranges))
+   if (++range > intel_pt_validate_hw_cap(PT_CAP_num_address_ranges))
        return -EOPNOTSUPP;
}

@@ -1507,12 +1510,12 @@ static __init int pt_init(void)
    if (ret)
        return ret;

-   if (!pt_cap_get(PT_CAP_topa_output)) {
+   if (!intel_pt_validate_hw_cap(PT_CAP_topa_output)) {
        pr_warn("ToPA output is not supported on this CPU\n");
        return -ENODEV;
    }

-   if (!pt_cap_get(PT_CAP_topa_multiple_entries))
+   if (!intel_pt_validate_hw_cap(PT_CAP_topa_multiple_entries))
        pt_pmu.pmu.capabilities =
            PERF_PMU_CAP_AUX_NO_SG | PERF_PMU_CAP_AUX_SW_DOUBLEBUF;

@@ -1530,7 +1533,7 @@ static __init int pt_init(void)
    pt_pmu.pmu.addr_filters_sync = pt_event_addr_filters_sync;
    pt_pmu.pmu.addr_filters_validate = pt_event_addr_filters_validate;
    pt_pmu.pmu.nr_addr_filters =
-   pt_cap_get(PT_CAP_num_address_ranges);
+   intel_pt_validate_hw_cap(PT_CAP_num_address_ranges);

    ret = perf_pmu_register(&pt_pmu.pmu, "intel_pt", -1);

```

**diff --git a/arch/x86/events/intel/pt.h b/arch/x86/events/intel/pt.h**

```

index 0050ca1..269e15a 100644
--- a/arch/x86/events/intel/pt.h
+++ b/arch/x86/events/intel/pt.h
@@ -45,30 +45,9 @@ struct topa_entry {
    u64      rsvd4      : 16;
};

-#define PT_CPUID_LEAVES                2
-#define PT_CPUID_REGS_NUM              4 /* number of registers (eax, ebx, ecx, edx) */
-
/* TSC to Core Crystal Clock Ratio */
#define CPUID_TSC_LEAF                  0x15

-enum pt_capabilities {
-    PT_CAP_max_subleaf = 0,
-    PT_CAP_cr3_filtering,
-    PT_CAP_psb_cyc,
-    PT_CAP_ip_filtering,
-    PT_CAP_mtc,
-    PT_CAP_ptwrite,
-    PT_CAP_power_event_trace,
-    PT_CAP_topa_output,
-    PT_CAP_topa_multiple_entries,
-    PT_CAP_single_range_output,
-    PT_CAP_payloads_lip,
-    PT_CAP_num_address_ranges,
-    PT_CAP_mtc_periods,
-    PT_CAP_cycle_thresholds,
-    PT_CAP_psb_periods,
-};
-
struct pt_pmu {
    struct pmu
        pmu;
    u32
        caps[PT_CPUID_REGS_NUM * PT_CPUID_LEAVES];
diff --git a/arch/x86/include/asm/intel_pt.h b/arch/x86/include/asm/intel_pt.h
index b523f51..fa4b4fd 100644
--- a/arch/x86/include/asm/intel_pt.h
+++ b/arch/x86/include/asm/intel_pt.h
@@ -2,10 +2,33 @@
-#ifndef _ASM_X86_INTEL_PT_H
-#define _ASM_X86_INTEL_PT_H
-
+#define PT_CPUID_LEAVES                2
+#define PT_CPUID_REGS_NUM              4 /* number of registers (eax, ebx, ecx, edx) */
+
+enum pt_capabilities {
+    PT_CAP_max_subleaf = 0,
+    PT_CAP_cr3_filtering,
+    PT_CAP_psb_cyc,
+    PT_CAP_ip_filtering,
+    PT_CAP_mtc,

```

```
+    PT_CAP_ptwrite,  
+    PT_CAP_power_event_trace,  
+    PT_CAP_topa_output,  
+    PT_CAP_topa_multiple_entries,  
+    PT_CAP_single_range_output,  
+    PT_CAP_payloads_lip,  
+    PT_CAP_num_address_ranges,  
+    PT_CAP_mtc_periods,  
+    PT_CAP_cycle_thresholds,  
+    PT_CAP_psb_periods,  
+};  
+  
#if defined(CONFIG_PERF_EVENTS) && defined(CONFIG_CPU_SUP_INTEL)  
void cpu_emergency_stop_pt(void);  
+extern u32 intel_pt_validate_hw_cap(enum pt_capabilities cap);  
#else  
static inline void cpu_emergency_stop_pt(void) {}  
+static inline u32 intel_pt_validate_hw_cap(enum pt_capabilities cap) { return 0; }  
#endif  
  
#endif /* _ASM_X86_INTEL_PT_H */
```

patchwork (<http://jk.ozlabs.org/projects/patchwork/>) patch tracking system | version v2.1.0 | [about patchwork \(/about/\)](#)