

[v13,07/12] KVM: x86: Add Intel Processor Trace cpuid emulation

10654367

[diff \(/patch/10654367/raw/\)](/patch/10654367/raw/)[mbox \(/patch/10654367/mbox/\)](/patch/10654367/mbox/)[series \(/series/34463/mbox/\)](/series/34463/mbox/)**Message ID** 1540368316-12998-8-git-send-email-luwei.kang@intel.com**State** New**Headers** show**Series** Intel Processor Trace virtualization enabling**Related** show

Commit Message

Kang, Luwei (/project/kvm/list/?submitter=168537)

Oct. 24, 2018, 8:05 a.m. UTC

From: Chao Peng <chao.p.peng@linux.intel.com>Expose Intel Processor Trace to guest only when
the PT works in Host-Guest mode.**Signed-off-by:** Chao Peng <chao.p.peng@linux.intel.com>**Signed-off-by:** Luwei Kang <luwei.kang@intel.com>

```

arch/x86/include/asm/kvm_host.h | 1 +
arch/x86/kvm/cpuid.c             | 22 ++++++
arch/x86/kvm/svm.c               | 6 +++++
arch/x86/kvm/vmx.c               | 6 +++++
4 files changed, 33 insertions(+), 2 deletions(-)

```

Patch

10654367

[diff \(/patch/10654367/raw/\)](/patch/10654367/raw/)[mbox \(/patch/10654367/mbox/\)](/patch/10654367/mbox/)[series \(/series/34463/mbox/\)](/series/34463/mbox/)

```

diff --git a/arch/x86/include/asm/kvm_host.h b/arch/x86/include/asm/kvm_host.h
index 55e51ff..9ab7ac0 100644
--- a/arch/x86/include/asm/kvm_host.h
+++ b/arch/x86/include/asm/kvm_host.h
@@ -1105,6 +1105,7 @@ struct kvm_x86_ops {
    bool (*mpx_supported)(void);
    bool (*xsaves_supported)(void);
    bool (*umip_emulated)(void);
+   bool (*pt_supported)(void);

    int (*check_nested_events)(struct kvm_vcpu *vcpu, bool external_intr);
    void (*request_immediate_exit)(struct kvm_vcpu *vcpu);
diff --git a/arch/x86/kvm/cpuid.c b/arch/x86/kvm/cpuid.c
index 7bcfa61..05b8fb4 100644
--- a/arch/x86/kvm/cpuid.c
+++ b/arch/x86/kvm/cpuid.c
@@ -337,6 +337,7 @@ static inline int __do_cpuid_ent(struct kvm_cpuid_entry2 *entry, u32 function,
    unsigned f_mpx = kvm_mpx_supported() ? F(MPX) : 0;
    unsigned f_xsaves = kvm_x86_ops->xsaves_supported() ? F(XSAVES) : 0;
    unsigned f_umip = kvm_x86_ops->umip_emulated() ? F(UMIP) : 0;
+   unsigned f_intel_pt = kvm_x86_ops->pt_supported() ? F(INTEL_PT) : 0;

    /* cpuid 1.edx */
    const u32 kvm_cpuid_1_edx_x86_features =
@@ -395,7 +396,7 @@ static inline int __do_cpuid_ent(struct kvm_cpuid_entry2 *entry, u32 function,
    F(BMI2) | F(ERMS) | f_invpaid | F(RTM) | f_mpx | F(RDSEED) |
    F(ADX) | F(SMAP) | F(AVX512IFMA) | F(AVX512F) | F(AVX512PF) |
    F(AVX512ER) | F(AVX512CD) | F(CLFLUSHOPT) | F(CLWB) | F(AVX512DQ) |
-   F(SHA_NI) | F(AVX512BW) | F(AVX512VL);
+   F(SHA_NI) | F(AVX512BW) | F(AVX512VL) | f_intel_pt;

    /* cpuid 0xD.1.eax */
    const u32 kvm_cpuid_D_1_eax_x86_features =
@@ -426,7 +427,7 @@ static inline int __do_cpuid_ent(struct kvm_cpuid_entry2 *entry, u32 function,
    switch (function) {
    case 0:
-       entry->eax = min(entry->eax, (u32)0xd);
+       entry->eax = min(entry->eax, (u32)(f_intel_pt ? 0x14 : 0xd));
        break;
    case 1:
        entry->edx &= kvm_cpuid_1_edx_x86_features;
@@ -603,6 +604,23 @@ static inline int __do_cpuid_ent(struct kvm_cpuid_entry2 *entry, u32 function,
    }
    break;
}

+ /* Intel PT */
+ case 0x14: {
+     int t, times = entry->eax;
+

```

```

+         if (!f_intel_pt)
+             break;
+
+         entry->flags |= KVM_CPUID_FLAG_SIGNIFCANT_INDEX;
+         for (t = 1; t <= times; ++t) {
+             if (*nent >= maxnent)
+                 goto out;
+             do_cpuid_1_ent(&entry[t], function, t);
+             entry[t].flags |= KVM_CPUID_FLAG_SIGNIFCANT_INDEX;
+             ++*nent;
+         }
+         break;
+     }
+     case KVM_CPUID_SIGNATURE: {
+         static const char signature[12] = "KVMKVMKVM\0\0";
+         const u32 *sigptr = (const u32 *)signature;
diff --git a/arch/x86/kvm/svm.c b/arch/x86/kvm/svm.c
index f416f5c7..6e8a61b 100644
--- a/arch/x86/kvm/svm.c
+++ b/arch/x86/kvm/svm.c
@@ -5904,6 +5904,11 @@ static bool svm_umip_emulated(void)
     return false;
 }

+static bool svm_pt_supported(void)
+{
+    return false;
+}
+
+static bool svm_has_wbinvd_exit(void)
+{
+    return true;
+}
@@ -7139,6 +7144,7 @@ static int nested_enable_evmcs(struct kvm_vcpu *vcpu,
     .mpx_supported = svm_mpx_supported,
     .xsaves_supported = svm_xsaves_supported,
     .umip_emulated = svm_umip_emulated,
+    .pt_supported = svm_pt_supported,

     .set_supported_cpuid = svm_set_supported_cpuid,

diff --git a/arch/x86/kvm/vmx.c b/arch/x86/kvm/vmx.c
index c4c4b76..692154c 100644
--- a/arch/x86/kvm/vmx.c
+++ b/arch/x86/kvm/vmx.c
@@ -11013,6 +11013,11 @@ static bool vmx_xsaves_supported(void)
     SECONDARY_EXEC_XSAVES;
 }

+static bool vmx_pt_supported(void)
+{
+    return (pt_mode == PT_MODE_HOST_GUEST);
+}

```

```
+}
+
static void vmx_recover_nmi_blocking(struct vcpu_vmx *vmx)
{
    u32 exit_intr_info;
@@ -15127,6 +15132,7 @@ static int vmx_set_nested_state(struct kvm_vcpu *vcpu,
    .mpx_supported = vmx_mpx_supported,
    .xsaves_supported = vmx_xsaves_supported,
    .umip_emulated = vmx_umip_emulated,
+    .pt_supported = vmx_pt_supported,

    .check_nested_events = vmx_check_nested_events,
    .request_immediate_exit = vmx_request_immediate_exit,
```

patchwork (<http://jk.ozlabs.org/projects/patchwork/>) patch tracking system | version v2.1.0 | [about patchwork \(/about/\)](/about/)