

[v13,11/12] KVM: x86: Set intercept for Intel PT MSR read/write

10654373

[diff \(/patch/10654373/raw/\)](/patch/10654373/raw/)[mbox \(/patch/10654373/mbox/\)](/patch/10654373/mbox/)[series \(/series/34463/mbox/\)](/series/34463/mbox/)**Message ID** 1540368316-12998-12-git-send-email-luwei.kang@intel.com**State** New**Headers** show**Series** Intel Processor Trace virtualization enabling**Related** show

Commit Message

Kang, Luwei (/project/kvm/list/?submitter=168537)

Oct. 24, 2018, 8:05 a.m. UTC

From: Chao Peng <chao.p.peng@linux.intel.com>

To save performance overhead, disable intercept Intel PT MSR read/write when Intel PT is enabled in guest.
MSR_IA32_RTIT_CTL is an exception that will always be intercepted.

Signed-off-by: Chao Peng <chao.p.peng@linux.intel.com>**Signed-off-by:** Luwei Kang <luwei.kang@intel.com>

```
arch/x86/kvm/vmx.c | 23 ++++++
1 file changed, 23 insertions(+)
```

Patch

10654373

[diff \(/patch/10654373/raw/\)](/patch/10654373/raw/)[mbox \(/patch/10654373/mbox/\)](/patch/10654373/mbox/)[series \(/series/34463/mbox/\)](/series/34463/mbox/)

```

diff --git a/arch/x86/kvm/vmx.c b/arch/x86/kvm/vmx.c
index a568d49..ed247dd 100644
--- a/arch/x86/kvm/vmx.c
+++ b/arch/x86/kvm/vmx.c
@@ -1333,6 +1333,7 @@ static bool nested_vmx_is_page_fault_vmexit(struct vmcs12 *vmcs12,
 static void vmx_update_msr_bitmap(struct kvm_vcpu *vcpu);
 static void __always_inline vmx_disable_intercept_for_msr(unsigned long *msr_bitmap,
                                                           u32 msr, int type);
+static void pt_set_intercept_for_msr(struct vcpu_vmx *vmx, bool flag);

 static DEFINE_PER_CPU(struct vmcs *, vmxarea);
 static DEFINE_PER_CPU(struct vmcs *, current_vmcs);
@@ -4558,6 +4559,7 @@ static int vmx_set_msr(struct kvm_vcpu *vcpu, struct msr_data *msr_info)
     vmx_rtit_ctl_check(vcpu, data)
     return 1;
     vmcs_write64(GUEST_IA32_RTIT_CTL, data);
+    pt_set_intercept_for_msr(vmx, !(data & RTIT_CTL_TRACEEN));
+    vmx->pt_desc.guest.ctl = data;
     break;
     case MSR_IA32_RTIT_STATUS:
@@ -6414,6 +6416,27 @@ static void vmx_update_msr_bitmap(struct kvm_vcpu *vcpu)
     vmx->msr_bitmap_mode = mode;
 }

+static void pt_set_intercept_for_msr(struct vcpu_vmx *vmx, bool flag)
+{
+    unsigned long *msr_bitmap = vmx->vmcs01.msr_bitmap;
+    u32 i;
+
+    vmx_set_intercept_for_msr(msr_bitmap, MSR_IA32_RTIT_STATUS,
+                              MSR_TYPE_RW, flag);
+    vmx_set_intercept_for_msr(msr_bitmap, MSR_IA32_RTIT_OUTPUT_BASE,
+                              MSR_TYPE_RW, flag);
+    vmx_set_intercept_for_msr(msr_bitmap, MSR_IA32_RTIT_OUTPUT_MASK,
+                              MSR_TYPE_RW, flag);
+    vmx_set_intercept_for_msr(msr_bitmap, MSR_IA32_RTIT_CR3_MATCH,
+                              MSR_TYPE_RW, flag);
+    for (i = 0; i < vmx->pt_desc.addr_range; i++) {
+        vmx_set_intercept_for_msr(msr_bitmap,
+                                  MSR_IA32_RTIT_ADDR0_A + i * 2, MSR_TYPE_RW, flag);
+        vmx_set_intercept_for_msr(msr_bitmap,
+                                  MSR_IA32_RTIT_ADDR0_B + i * 2, MSR_TYPE_RW, flag);
+    }
+}
+
 static bool vmx_get_enable_apicv(struct kvm_vcpu *vcpu)
 {
     return enable_apicv;
 }

```

patchwork (<http://jk.ozlabs.org/projects/patchwork/>) patch tracking system | version v2.1.0 | [about patchwork \(/about/\)](#)